

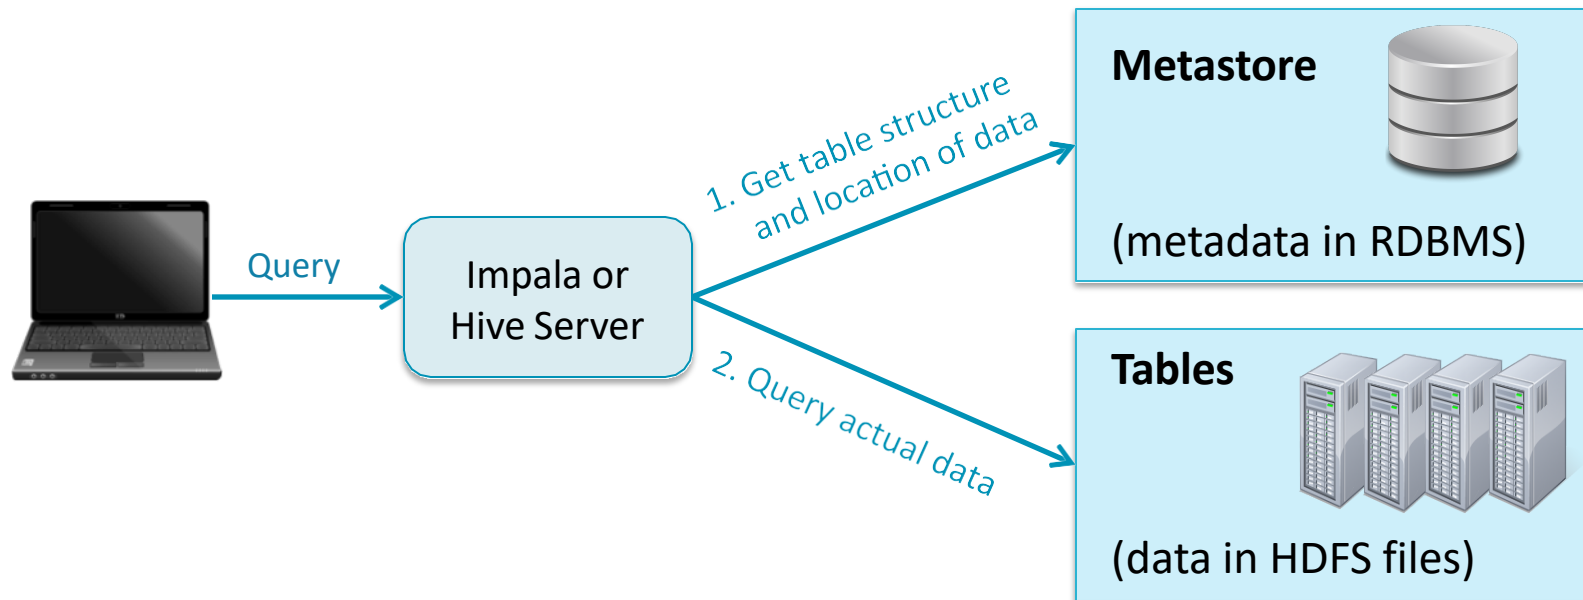
# Apache Hive Data Management

---



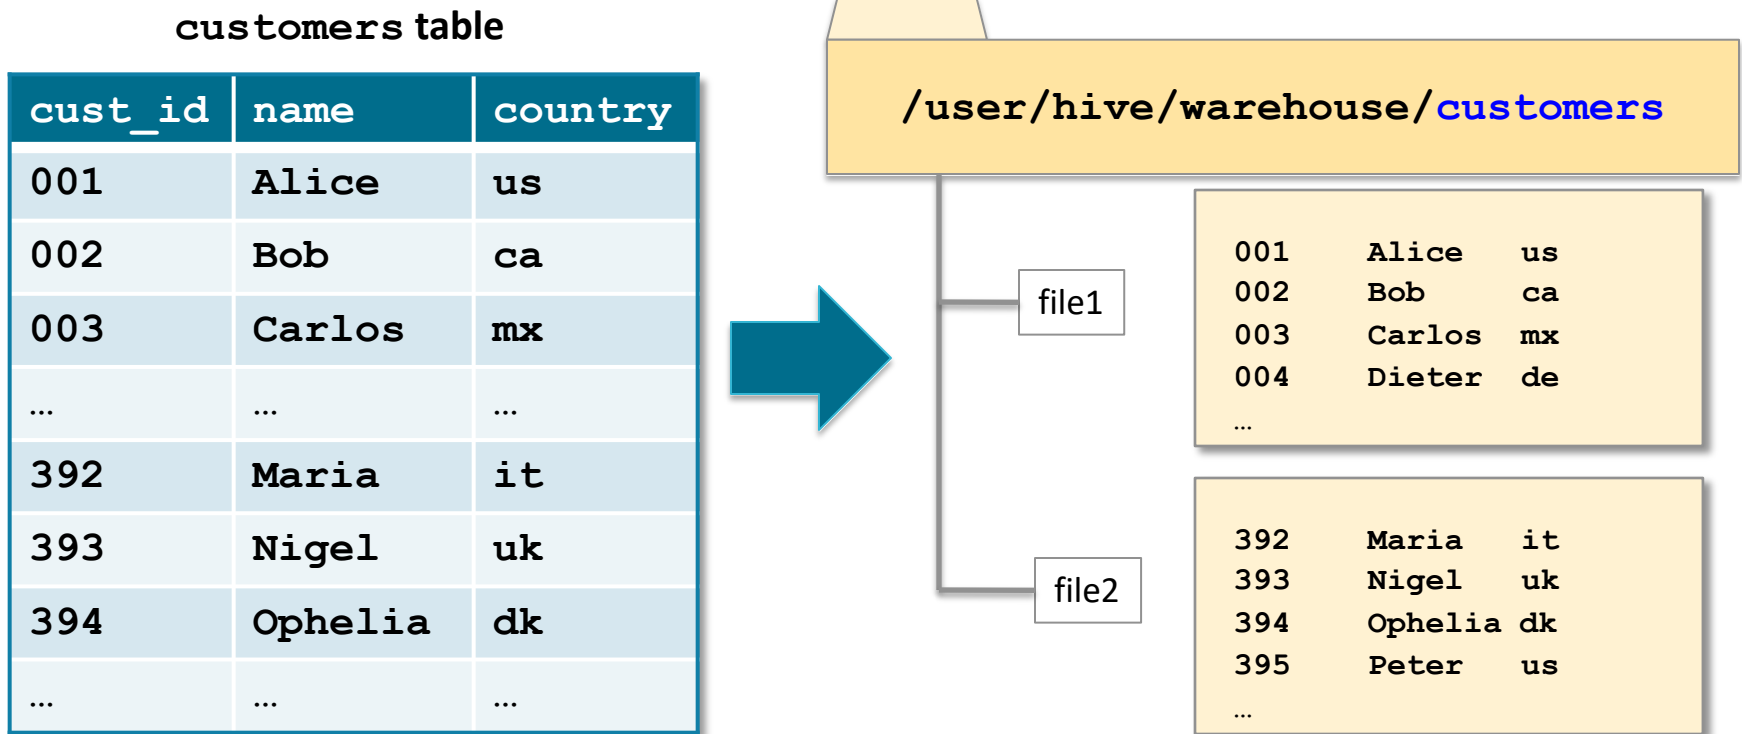
## Recap: How Hive and Impala Load and Store Data

- **Hive and Impala use the metastore to determine data format and location**
  - The query itself operates on data stored in a filesystem (typically HDFS)



# The Warehouse Directory

- By default, Hive and Impala store data in the HDFS directory `/user/hive/warehouse`
- Each table is a subdirectory containing any number of files



# Creating a Database

---

- **Hive and Impala databases are simply namespaces**

- Helps to organize your tables

- **To create a new database**

```
CREATE DATABASE dualcore;
```

1. Adds the database definition to the metastore
2. Creates a storage directory in HDFS

For example, `/user/hive/warehouse/dualcore.db`

- **To conditionally create a new database**

- Avoids error in case database already exists (useful for scripting)

```
CREATE DATABASE IF NOT EXISTS dualcore;
```

## Creating a Table (1)

---

- Basic syntax for creating a table:

```
CREATE TABLE dbname.tablename (colname DATATYPE, ...)  
  ROW FORMAT DELIMITED  
    FIELDS TERMINATED BY char  
  STORED AS {TEXTFILE|SEQUENCEFILE|...};
```

- Creates a subdirectory in the database's warehouse directory in HDFS

- Default database:

*/user/hive/warehouse/tablename*

- Named database:

*/user/hive/warehouse/dbname.db/tablename*

## Example Table Definition

---

- The following example creates a new table named `jobs`
  - Data stored as text with four comma-separated fields per line

```
CREATE TABLE jobs (  
    id INT,  
    title STRING,  
    salary INT,  
    posted TIMESTAMP  
)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ',';
```

- Example of corresponding record for the table above

```
1,Data Analyst,135000,2016-12-21 15:52:03
```

# Controlling Table Data Location

---

- By default, table data is stored in the warehouse directory
- This is not always ideal
  - Data might be part of a bigger workflow
- Use **LOCATION** to specify the directory where table data resides

```
CREATE TABLE jobs (  
    id INT,  
    title STRING,  
    salary INT,  
    posted TIMESTAMP  
)  
ROW FORMAT DELIMITED  
    FIELDS TERMINATED BY ','  
LOCATION '/dualcore/jobs';
```

# Externally Managed Tables

---

- **CAUTION: Dropping a table removes its data in HDFS**
- **Using EXTERNAL when creating the table avoids this behavior**
  - Dropping an *external (unmanaged)* table removes only its *metadata*

```
CREATE EXTERNAL TABLE adclicks (  
    campaign_id STRING,  
    click_time TIMESTAMP,  
    keyword STRING,  
    site STRING,  
    placement STRING,  
    was_clicked BOOLEAN,  
    cost SMALLINT  
)  
LOCATION '/dualcore/ad_data';
```



# Data Validation

---

- **Hive and Impala are *schema-on-read***
  - Unlike an RDBMS, they do not validate data on insert
    - Files are simply moved into place
  - Loading data into tables is therefore very fast
  - Errors in file format will be discovered when queries are performed
- **Missing or invalid data will be represented as NULL**

## Loading Data from HDFS Files

---

- **To load data, simply add files to the table's directory in HDFS**

- Can be done directly using the **hdfs dfs** commands
- This example loads data from HDFS into the **sales** table

```
$ hdfs dfs -mv sales.txt /user/hive/warehouse/sales/
```

- **Alternatively, use the `LOAD DATA INPATH` command**

- Done from within Hive or Impala
- This *moves* data within HDFS, just like the command above
- Source can be either a file or directory

```
LOAD DATA INPATH '/incoming/etl/sales.txt'  
  INTO TABLE sales;
```

## Overwriting Data from Files

---

- **Add the OVERWRITE keyword to delete all records before import**
  - Removes all files within the table's directory
  - Then moves the new files into that directory

```
LOAD DATA INPATH '/incoming/etl/sales.txt'  
  OVERWRITE INTO TABLE sales;
```

## Loading Data from a Relational Database

---

- Sqoop has built-in support for importing data into Hive and Impala
- Add the `--hive-import` option to your Sqoop command
  - Creates the table in the metastore
  - Imports data from the RDBMS to the table's directory in HDFS

```
$ sqoop import \  
  --connect jdbc:mysql://localhost/dualcore \  
  --username training \  
  --password training \  
  --fields-terminated-by '\t' \  
  --table employees \  
  --hive-import \  
  --hive-database default \  
  --hive-table employees
```

# Removing a Database

---

- Removing a database is similar to creating it

```
DROP DATABASE dualcore;
```

```
DROP DATABASE IF EXISTS dualcore;
```

- These commands will fail if the database contains tables

- Add the **CASCADE** keyword to force removal
  - Supported in all production versions of Hive
  - Supported in Impala 2.3 (CDH 5.5.0) and higher

```
DROP DATABASE dualcore CASCADE;
```

**CAUTION:**  
This command might  
remove data in HDFS!

# Removing a Table

---

- **Table removal syntax is similar to database removal**

```
DROP TABLE customers;
```

```
DROP TABLE IF EXISTS customers;
```

- **Managed (internal) tables**
  - Metadata is removed
  - Data in HDFS is removed
  - *Caution: No rollback or undo feature!*
- **Unmanaged (external) tables**
  - Metadata is removed
  - Data in HDFS is *not* removed

## Creating Views

---

- **Views are conceptually like a table, but backed by a query**
  - You cannot directly add data to a view

```
CREATE VIEW order_info AS
  SELECT o.order_id, order_date, p.prod_id, brand, name
  FROM orders o
  JOIN order_details d
  ON (o.order_id = d.order_id)
  JOIN products p
  ON (d.prod_id = p.prod_id);
```

- **The query is now greatly simplified**

```
SELECT * FROM order_info WHERE order_id=6584288;
```

## Exploring Views

---

- **SHOW TABLES** lists the tables *and views* in a database
  - There is no separate command to list only views

```
SHOW TABLES;
```

- Use **DESCRIBE FORMATTED** to see a view's underlying query

```
DESCRIBE FORMATTED order_info;
```

- Use **SHOW CREATE TABLE** to display a statement to create the view

```
SHOW CREATE TABLE order_info;
```



# Modifying and Removing Views

---

- Use **ALTER VIEW** to change the underlying query

```
ALTER VIEW order_info AS  
    SELECT order_id, order_date FROM orders;
```

- Or to rename a view

```
ALTER VIEW order_info  
    RENAME TO order_information;
```

- Use **DROP VIEW** to remove a view

```
DROP VIEW order_info;
```

## Saving Query Output to a Table

---

- **SELECT** statements display their results on screen
- **To save results to a table, use INSERT OVERWRITE TABLE**
  - Destination table must already exist
  - Existing contents will be deleted

```
INSERT OVERWRITE TABLE nyc_customers
  SELECT * FROM customers
  WHERE state = 'NY' AND city = 'New York';
```

- **INSERT INTO TABLE** adds records without first deleting existing data

```
INSERT INTO TABLE nyc_customers
  SELECT * FROM customers
  WHERE state = 'NY' AND city = 'Brooklyn';
```



## Writing Output to HDFS in Hive

- **Hive also lets you save output to a directory in HDFS**
  - *Caution: Hive does not delete existing contents of the directory!*

```
INSERT OVERWRITE DIRECTORY '/dualcore/ny/'  
  ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t'  
  SELECT * FROM customers  
  WHERE state = 'NY';
```