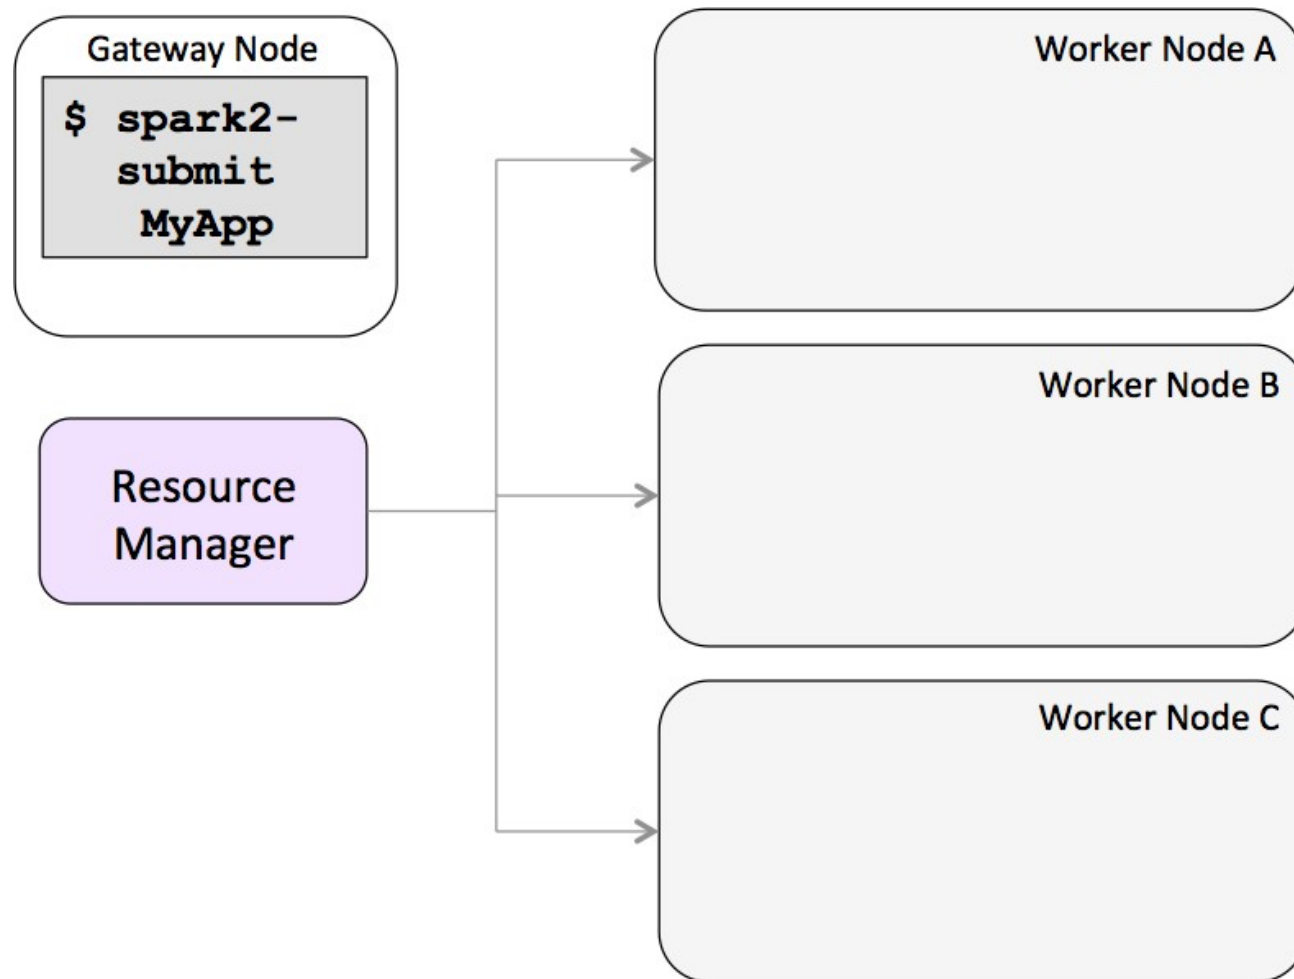
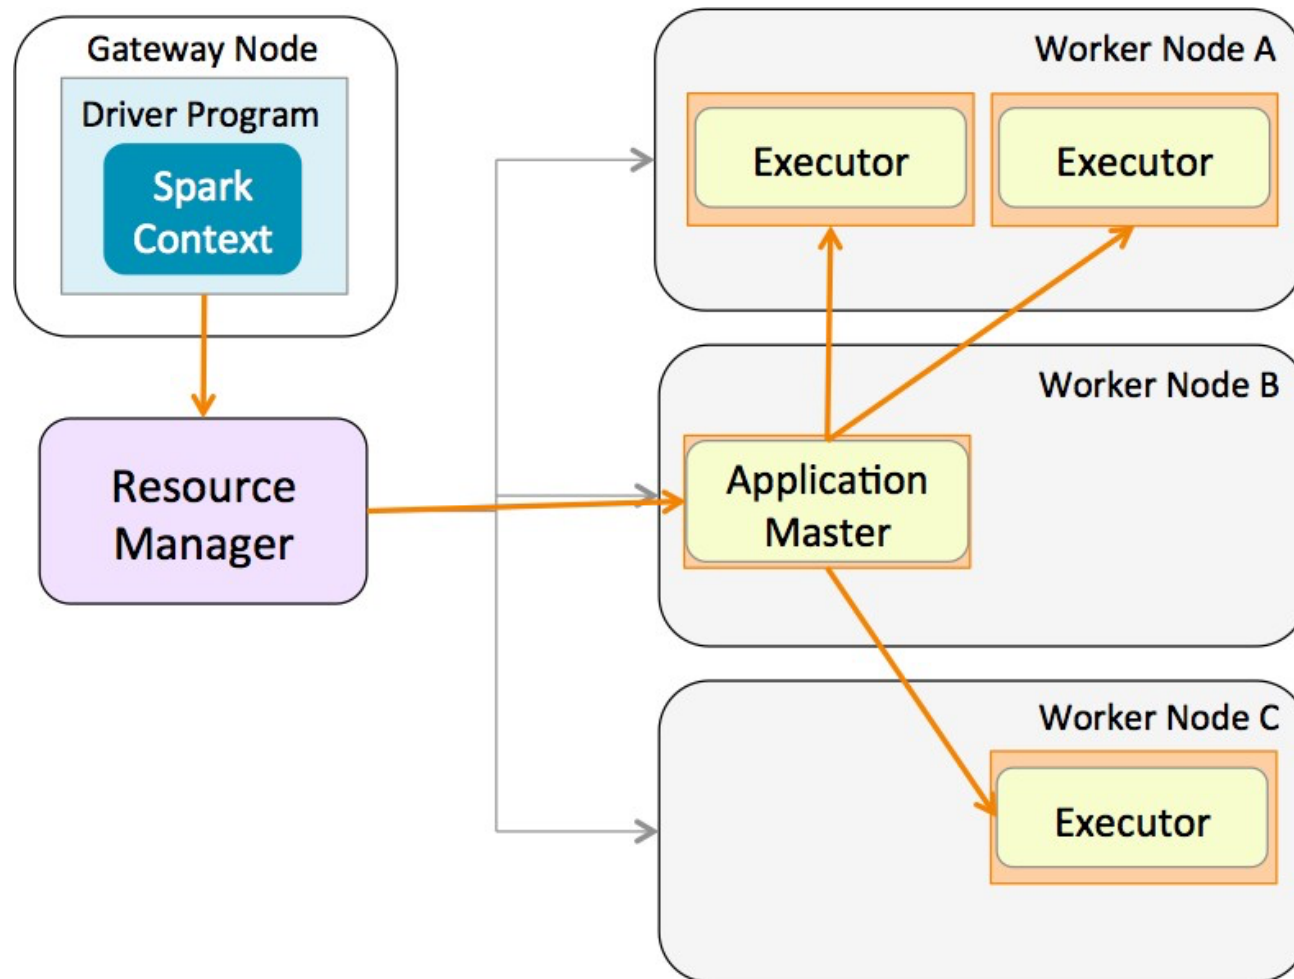


Distributed Processing

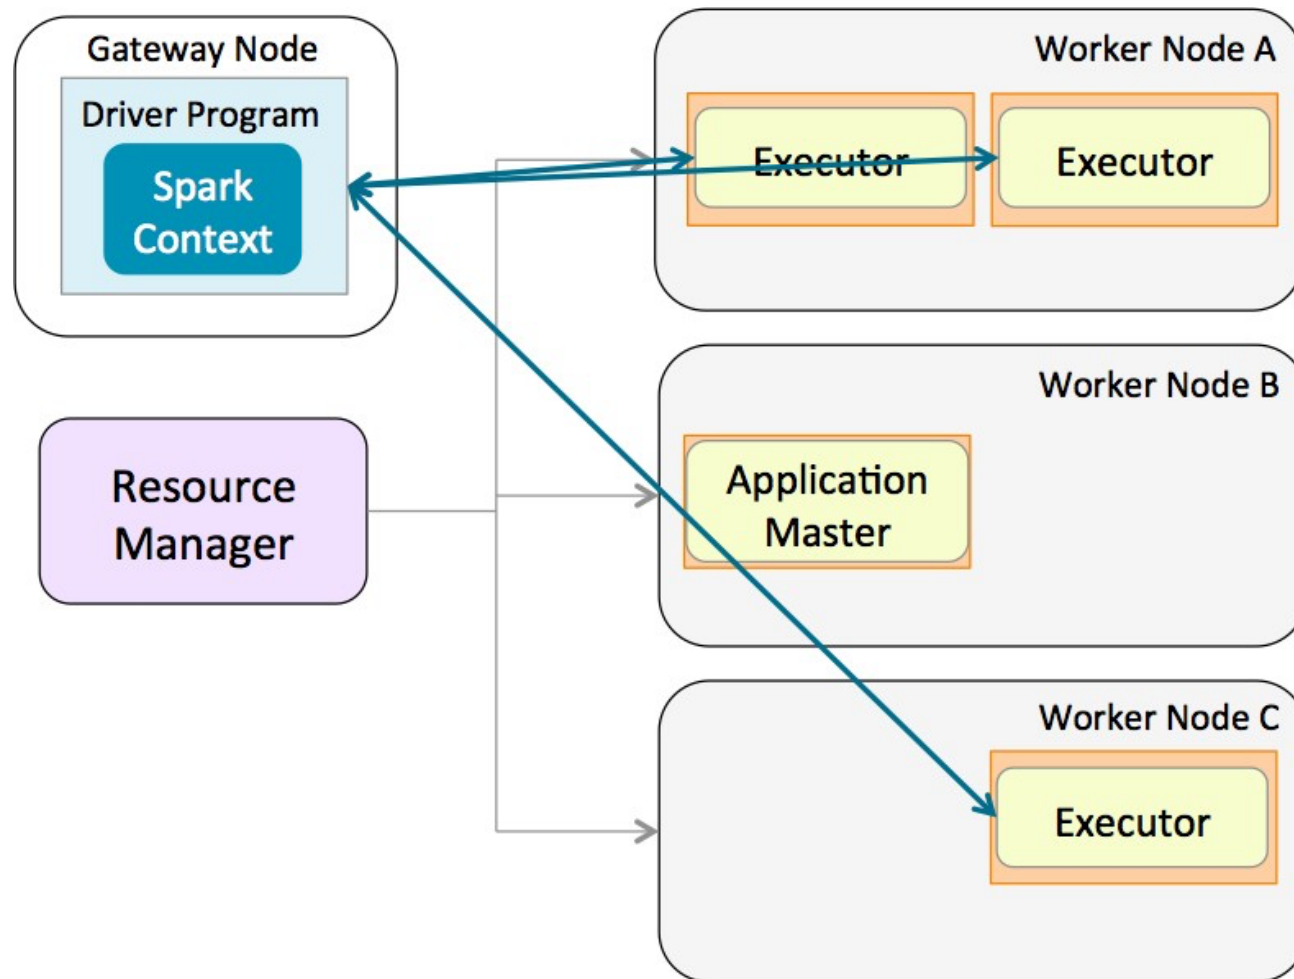
Review of Spark on YARN (1)



Review of Spark on YARN (2)




Review of Spark on YARN (3)



Data Partitioning (1)

- **Data in Datasets and DataFrames is managed by underlying RDDs**
- **Data in an RDD is *partitioned* across executors**
 - This is what makes RDDs *distributed*
 - Spark assigns tasks to process a partition to the executor managing that partition
- **Data Partitioning is done automatically by Spark**
 - In some cases, you can control how many partitions are created
 - More partitions = more parallelism



Data Partitioning (2)

- **Spark determines how to partition data in an RDD, Dataset, or DataFrame when**
 - The data source is read
 - An operation is performed on a DataFrame, Dataset, or RDD
 - Spark optimizes a query
 - You call `repartition` or `coalesce`

Partitioning from Data in Files

- **Partitions are determined when files are read**
 - Core Spark determines RDD partitioning based on location, number, and size of files
 - Usually each file is loaded into a single partition
 - Very large files are split across multiple partitions
 - Catalyst optimizer manages partitioning of RDDs that implement DataFrames and Datasets

Finding the Number of Partitions in an RDD

- You can view the number of partitions in an RDD by calling the function `getNumPartitions`

```
myRDD.getNumPartitions
```

Language: *Scala*


```
myRDD.getNumPartitions()
```

Language: *Python*

Example: Average Word Length by Letter (1)

```
avglens = sc.textFile(mydata)
```


Language: *Python*



Example: Average Word Length by Letter (2)

```
avglens = sc.textFile(mydata) \  
    .flatMap(lambda line: line.split(' '))
```

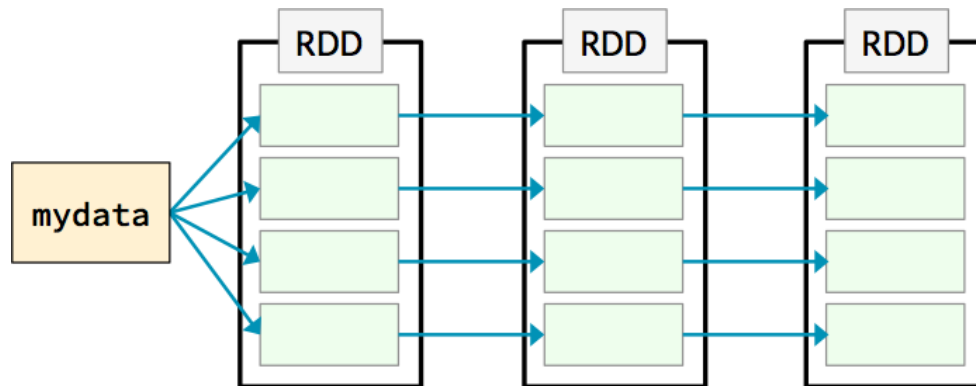
Language: *Python*



Example: Average Word Length by Letter (3)

```
avglens = sc.textFile(mydata) \  
    .flatMap(lambda line: line.split(' ')) \  
    .map(lambda word: (word[0], len(word)))
```

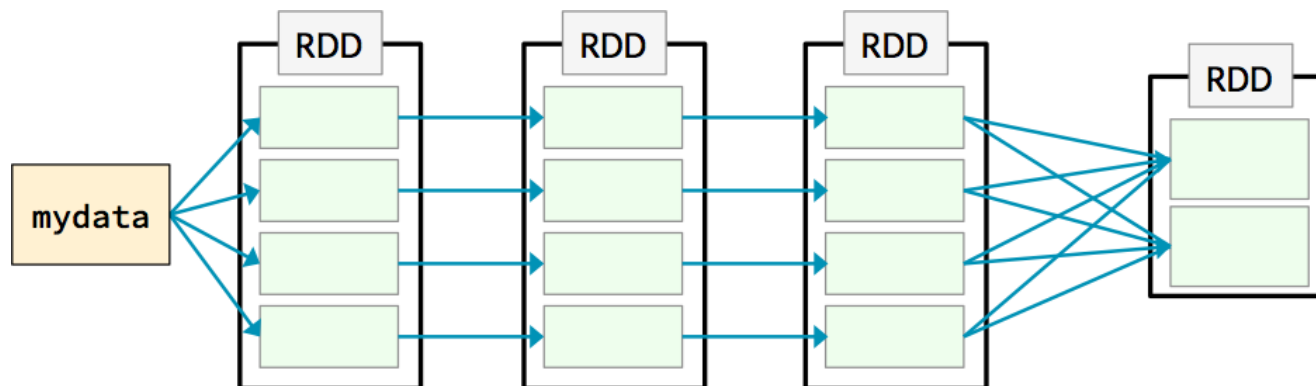
Language: *Python*



Example: Average Word Length by Letter (4)

```
avglens = sc.textFile(mydata) \  
    .flatMap(lambda line: line.split(' ')) \  
    .map(lambda word: (word[0],len(word))) \  
    .groupByKey()
```

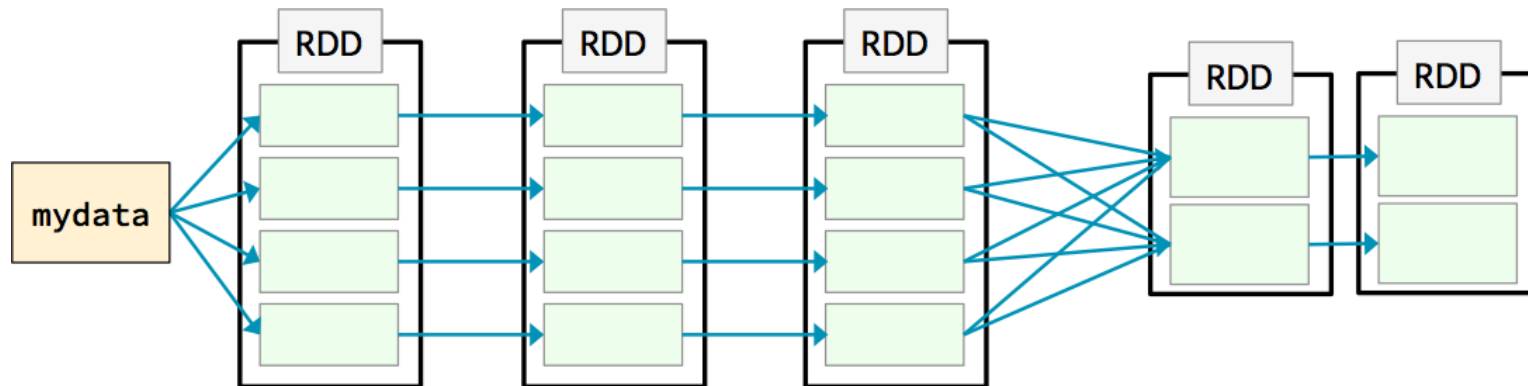
Language: *Python*



Example: Average Word Length by Letter (5)

```
avglens = sc.textFile(mydata) \  
    .flatMap(lambda line: line.split(' ')) \  
    .map(lambda word: (word[0],len(word))) \  
    .groupByKey() \  
    .map(lambda (k, values): \  
        (k, sum(values)/len(values)))
```

Language: Python



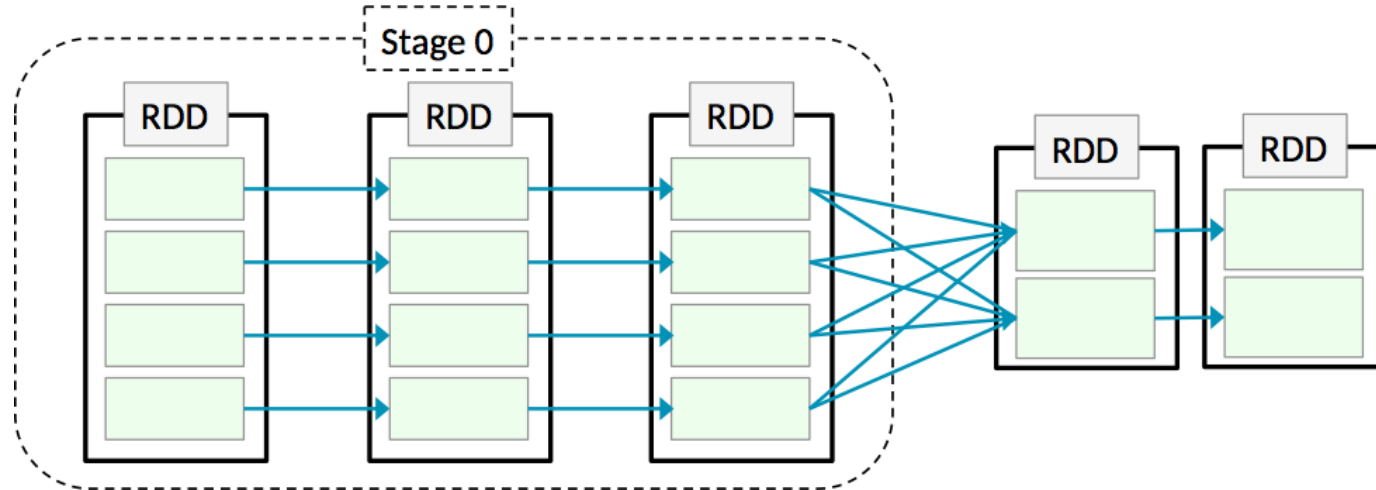
Stages and Tasks

- A **task** is a series of operations that work on the same partition and are pipelined together
- **Stages** group together tasks that can run in parallel on different partitions of the same RDD
- **Jobs** consist of all the stages that make up a query
- **Catalyst optimizes partitions and stages when using DataFrames and Datasets**
 - Core Spark provides limited optimizations when you work directly with RDDs
 - You need to code most RDD optimizations manually
 - To improve performance, be aware of how tasks and stages are executed when working with RDDs

Example: Query Stages and Tasks (1)

```
avglens = sc.textFile(mydata) \  
    .flatMap(lambda line: line.split(' ')) \  
    .map(lambda word: (word[0], len(word))) \  
    .groupByKey() \  
    .map(lambda (k, values): \  
        (k, sum(values)/len(values)))  
  
avglens.saveAsTextFile("avglen-output")
```

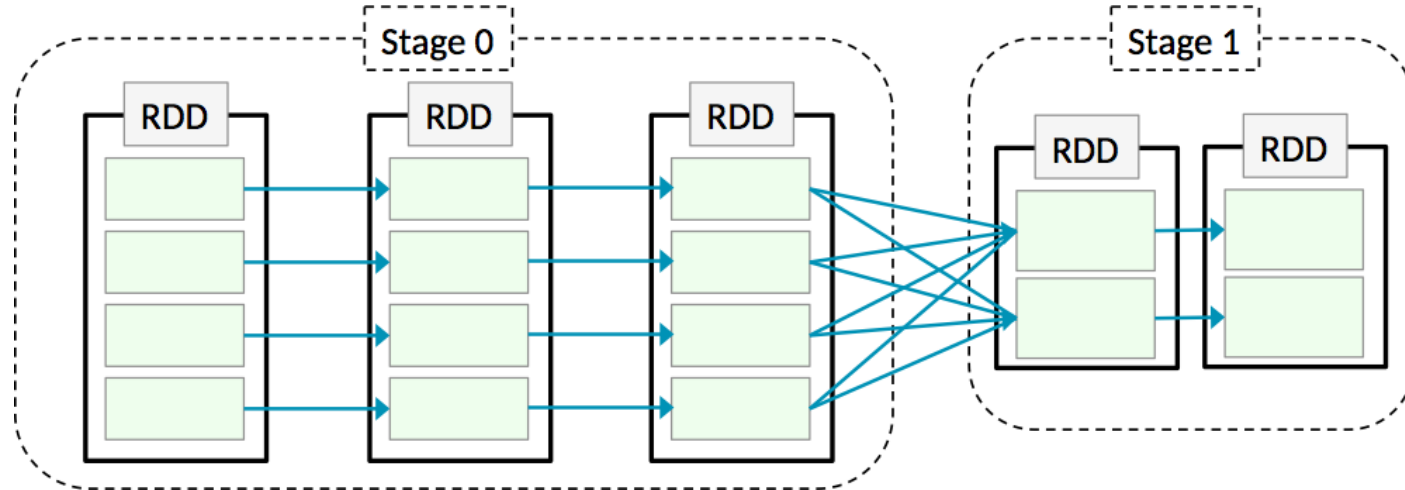
Language: *Python*



Example: Query Stages and Tasks (2)

```
avglens = sc.textFile(mydata) \  
    .flatMap(lambda line: line.split(' ')) \  
    .map(lambda word: (word[0], len(word))) \  
    .groupByKey() \  
    .map(lambda (k, values): \  
        (k, sum(values)/len(values)))  
  
avglens.saveAsTextFile("avglen-output")
```

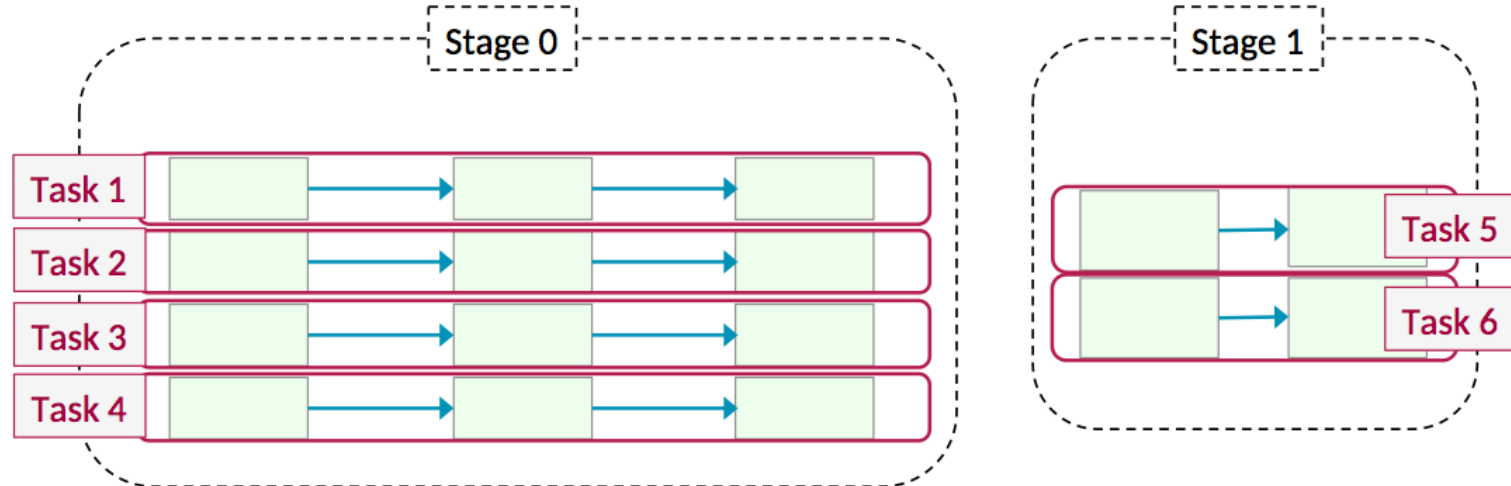
Language: *Python*



Example: Query Stages and Tasks (3)

```
avglens = sc.textFile(mydata) \  
    .flatMap(lambda line: line.split(' ')) \  
    .map(lambda word: (word[0],len(word))) \  
    .groupByKey() \  
    .map(lambda (k, values): \  
        (k, sum(values)/len(values)))  
  
avglens.saveAsTextFile("avglen-output")
```

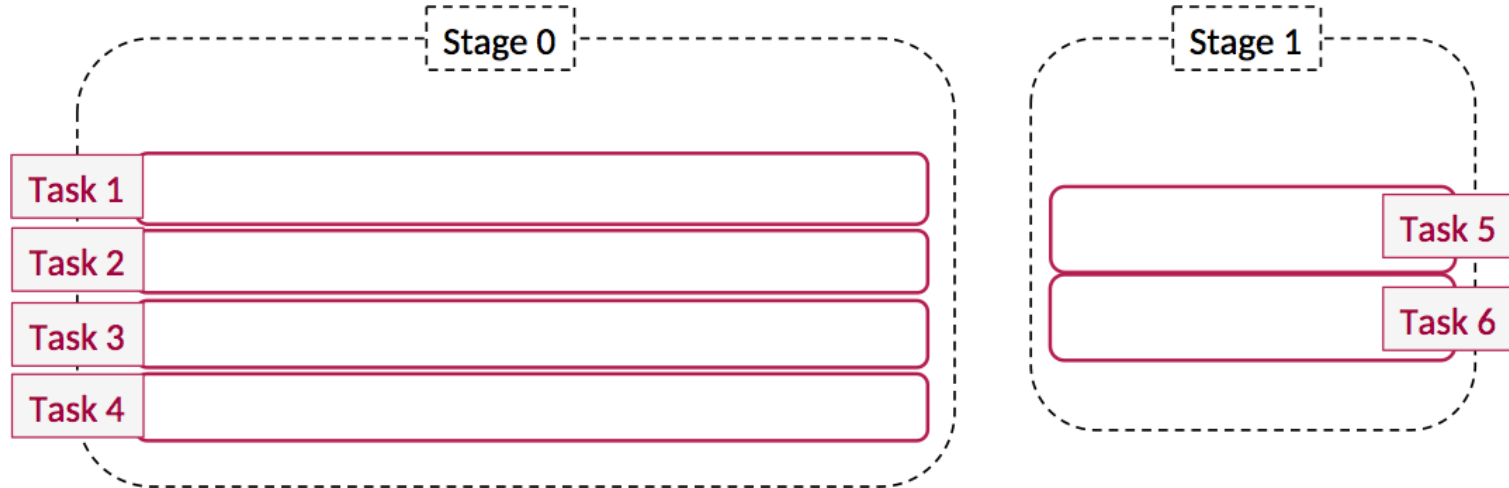
Language: Python



Example: Query Stages and Tasks (4)

```
avglens = sc.textFile(mydata) \  
    .flatMap(lambda line: line.split(' ')) \  
    .map(lambda word: (word[0],len(word))) \  
    .groupByKey() \  
    .map(lambda (k, values): \  
        (k, sum(values)/len(values)))  
  
avglens.saveAsTextFile("avglen-output")
```

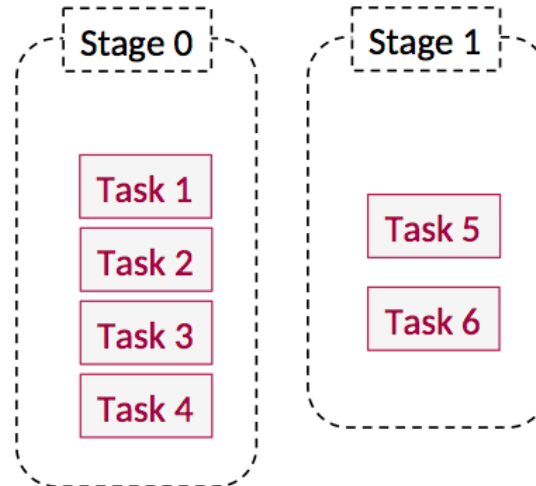
Language: *Python*



Example: Query Stages and Tasks (5)

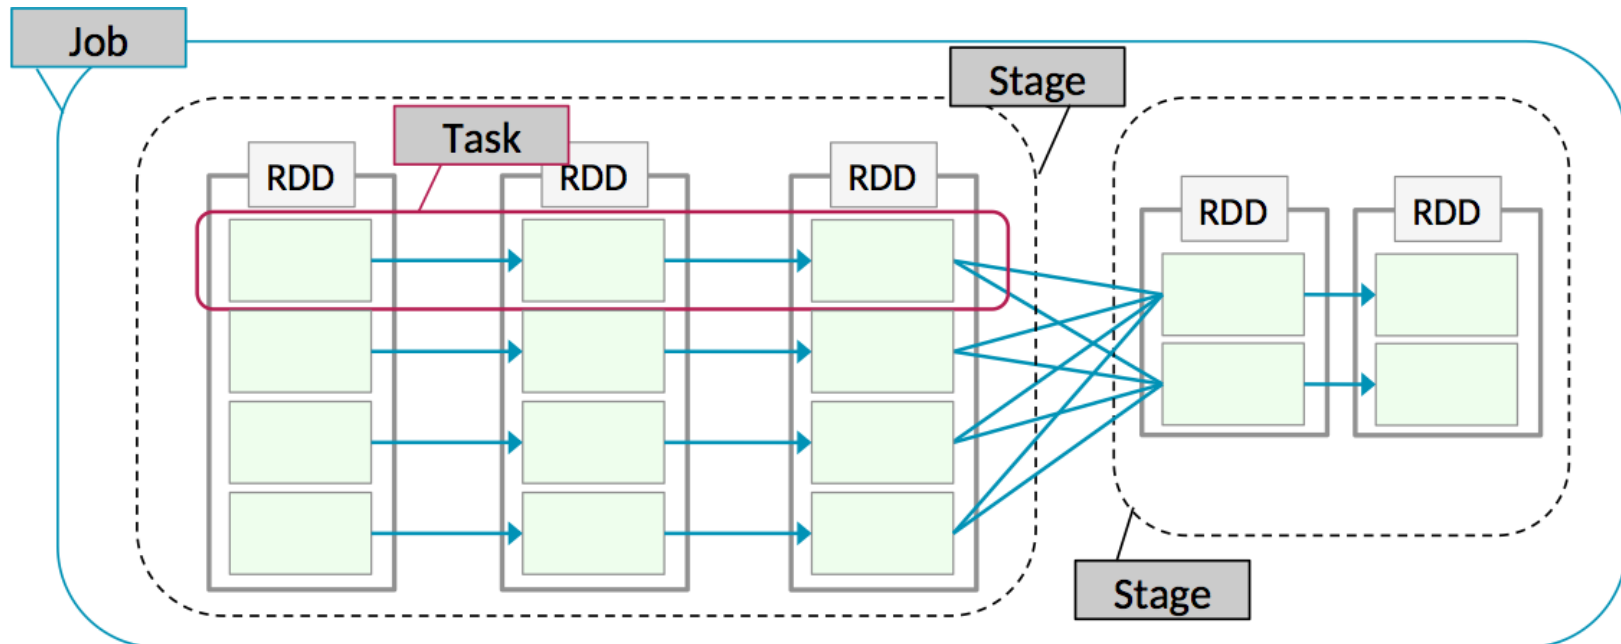
```
avglens = sc.textFile(mydata) \  
    .flatMap(lambda line: line.split(' ')) \  
    .map(lambda word: (word[0], len(word))) \  
    .groupByKey() \  
    .map(lambda (k, values): \  
        (k, sum(values)/len(values)))  
  
avglens.saveAsTextFile("avglen-output")
```

Language: *Python*



Summary of Spark Terminology

- **Job**—a set of tasks executed as a result of an *action*
- **Stage**—a set of tasks in a job that can be executed in parallel
- **Task**—an individual unit of work sent to one executor
- **Application**—the set of jobs managed by a single driver





Execution Plans

- **Spark creates an execution plan for each job in an application**
- **Catalyst creates SQL, Dataset, and DataFrame execution plans**
 - Highly optimized
- **Core Spark creates execution plans for RDDs**
 - Based on RDD lineage
 - Limited optimization

How Execution Plans are Created

- Spark constructs a **DAG (Directed Acyclic Graph)** based on **RDD dependencies**
- **Narrow dependencies**
 - Each partition in the child RDD depends on just one partition of the parent RDD
 - No shuffle required between executors
 - Can be pipelined into a single stage
 - Examples: `map`, `filter`, and `union`
- **Wide (or *shuffle*) dependencies**
 - Child partitions depend on multiple partitions in the parent RDD
 - Defines a new stage
 - Examples: `reduceByKey`, `join`, and `groupByKey`



Controlling the Number of Partitions in RDDs (1)

- **Partitioning determines how queries execute on a cluster**
 - More partitions = more parallel tasks
 - Cluster will be under-utilized if there are too few partitions
 - But too many partitions will increase overhead without an offsetting increase in performance
- **Catalyst controls partitioning for SQL, DataFrame, and Dataset queries**
- **You can control how many partitions are created for RDD queries**

Controlling the Number of Partitions in RDDs (2)

- **Specify the number of partitions when data is read**
 - Default partitioning is based on size and number of the files (minimum is two)
 - Specify a different minimum number when reading a file

```
myRDD = sc.textFile(myfile, 5)
```

- **Manually repartition**
 - Create a new RDD with a specified number of partitions using `repartition` or `coalesce`
 - `coalesce` reduces the number of partitions without requiring a shuffle
 - `repartition` shuffles the data into more or fewer partitions

```
newRDD = myRDD.repartition(15)
```


Controlling the Number of Partitions in RDDs (3)

- **Specify the number of partitions created by transformations**
 - Wide (shuffle) operations such as `reduceByKey` and `join` repartition data
 - By default, the number of partitions created is based on the number of partitions of the parent RDD(s)
 - Choose a different default by configuring the `spark.default.parallelism` property

```
spark.default.parallelism 15
```

- Override the default with the optional `numPartitions` operation parameter

```
countRDD = wordsRDD. \  
    reduceByKey(lambda v1, v2: v1 + v2, 15)
```

Catalyst Optimizer

- **Catalyst can improve SQL, DataFrame, and Dataset query performance by optimizing the DAG to**
 - Minimize data transfer between executors
 - Such as *broadcast* joins—small data sets are pushed to the executors where the larger data sets reside
 - Minimize wide (shuffle) operations
 - Such as unioning two RDDs—grouping, sorting, and joining do not require shuffling
 - Pipeline as many operations into a single stage as possible
 - Generate code for a whole stage at run time
 - Break a query job into multiple jobs, executed in a series

Catalyst Execution Plans

- **Execution plans for DataFrame, Dataset, and SQL queries include the following phases**
 - **Parsed logical plan**—calculated directly from the sequence of operations specified in the query
 - **Analyzed logical plan**—resolves relationships between data sources and columns
 - **Optimized logical plan**—applies rule-based optimizations
 - **Physical plan**—describes the actual sequence of operations
 - **Code generation**—generates bytecode to run on each node, based on a cost model

Viewing Catalyst Execution Plans

- **You can view SQL, DataFrame, and Dataset (Catalyst) execution plans**
 - Use DataFrame/Dataset `explain`
 - Shows only the physical execution plan by default
 - Pass `true` to see the full execution plan
 - Use **SQL** tab in the Spark UI or history server
 - Shows details of execution after job runs

Example: Catalyst Execution Plan (1)

```
peopleDF = spark.read. \  
    option("header","true").csv("people.csv")  
pcodesDF = spark.read. \  
    option("header","true").csv("pcodes.csv")  
joinedDF = peopleDF.join(pcodesDF, "pcode")  
joinedDF.explain(True)
```

== Parsed Logical Plan ==

```
'Join UsingJoin(Inner,ArrayBuffer('pcode'))  
:- Relation[pcode#0,lastName#1,firstName#2,age#3] csv  
+- Relation[pcode#9,city#10,state#11] csv
```

Language: Python
continued on next slide...

Example: Catalyst Execution Plan (2)

== Analyzed Logical Plan ==

```
pcode: string, lastName: string, firstName: string, age:
  string, city: string, state: string
Project [pcode#0, lastName#1, firstName#2, age#3, city#10,
  state#11]
+- Join Inner, (pcode#0 = pcode#9)
   :- Relation[pcode#0,lastName#1,firstName#2,age#3] csv
   +- Relation[pcode#9,city#10,state#11] csv
```

== Optimized Logical Plan ==

```
Project [pcode#0, lastName#1, firstName#2, age#3, city#10,
  state#11]
+- Join Inner, (pcode#0 = pcode#9)
   :- Filter isnotnull(pcode#0)
   :   +- Relation[pcode#0,lastName#1,firstName#2,age#3] csv
   +- Filter isnotnull(pcode#9)
      +- Relation[pcode#9,city#10,state#11] csv
```

Language: Python
continued on next slide...

Example: Catalyst Execution Plan (3)

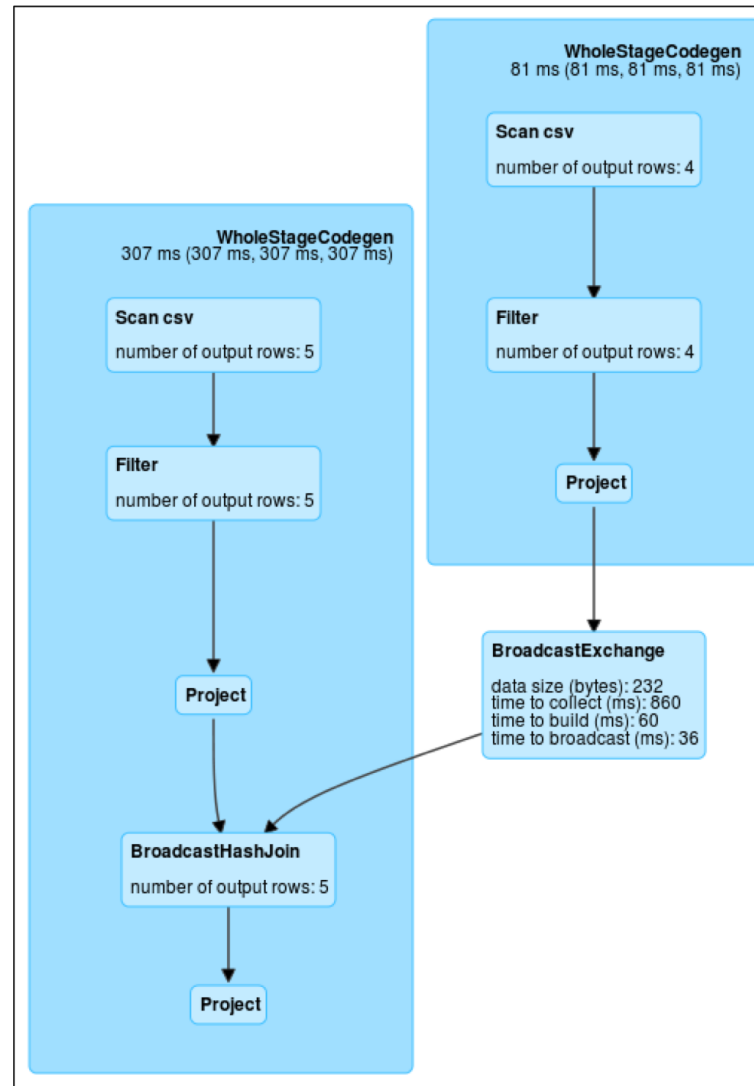
```
== Physical Plan ==
*Project [pcode#0, lastName#1, firstName#2, age#3, city#10,
state#11]
+- *BroadcastHashJoin [pcode#0], [pcode#9], Inner, BuildRight
  :- *Project [pcode#0, lastName#1, firstName#2, age#3]
  :   +- *Filter isnotnull(pcode#0)
  :     +- *Scan csv [pcode#0,lastName#1,firstName#2,age#3]
  Format: CSV, InputPaths: file:/home/training/people.csv,
  PushedFilters: [IsNotNull(pcode)], ReadSchema:
  struct<pcode:string,lastName:string,firstName:string,age:string>
  +- BroadcastExchange
  HashedRelationBroadcastMode(List(input[0, string, true]))
    +- *Project [pcode#9, city#10, state#11]
    +- *Filter isnotnull(pcode#9)
      +- *Scan csv [pcode#9,city#10,state#11]
  Format: CSV, InputPaths: file:/home/training/pcodes.csv,
  PushedFilters: [IsNotNull(pcode)], ReadSchema:
  struct<pcode:string,city:string,state:string></
pcode:string,city:string,state:string>
```

Language: Python

Example: Catalyst Execution Plan (4)

| Jobs Stages Storage Environment Executors SQL | | | | | |
|---|---|---------------------|----------|------|----|
| SQL | | | | | |
| Completed Queries | | | | | |
| ID | Description | Submitted | Duration | Jobs | |
| 0 | collect at <ipython-input-65-31b1b37d0504>:1 <small>details</small> | 2017/05/24 08:45:13 | 1 s | 31 | 32 |

Example: Catalyst Execution Plan (5)



Viewing RDD Execution Plans

- **You can view RDD (lineage-based) execution plans**
 - Use the RDD `toDebugString` function
 - Use **Jobs** and **Stages** tabs in the Spark UI or history server
 - Shows details of execution after job runs
- **Note that plans may be different depending on programming language**
 - Plan optimization rules vary

Example: RDD Execution Plan (1)

```
val peopleRDD = sc.textFile("people2.csv").keyBy(s => s.split(',') (0))
val pcodesRDD = sc.textFile("pcodes2.csv").keyBy(s => s.split(',') (0))
val joinedRDD = peopleRDD.join(pcodesRDD)
```

```
joinedRDD.toDebugString
```

```
(2) MapPartitionsRDD[8] at join at ... ①
  | MapPartitionsRDD[7] at join at ...
  | CoGroupedRDD[6] at join at ...
+- (2) MapPartitionsRDD[2] at keyBy at ... ②
  | | people2.csv MapPartitionsRDD[1] at textFile at ...
  | |
  | | people2.csv HadoopRDD[0] at ...
+- (2) MapPartitionsRDD[5] at keyBy at ... ③
  | pcodes2.csv MapPartitionsRDD[4] at textFile at ...
  |
  | ... ④
  | pcodes2.csv HadoopRDD[3] at textFile at ...
```

Language: *Scala*

① Stage 2

② Stage 1

③ Stage 0

④ Indents indicate stages (shuffle boundaries)

Example: RDD Execution Plan (2)

Jobs

Stages

Storage

Environment

Executors

SQL

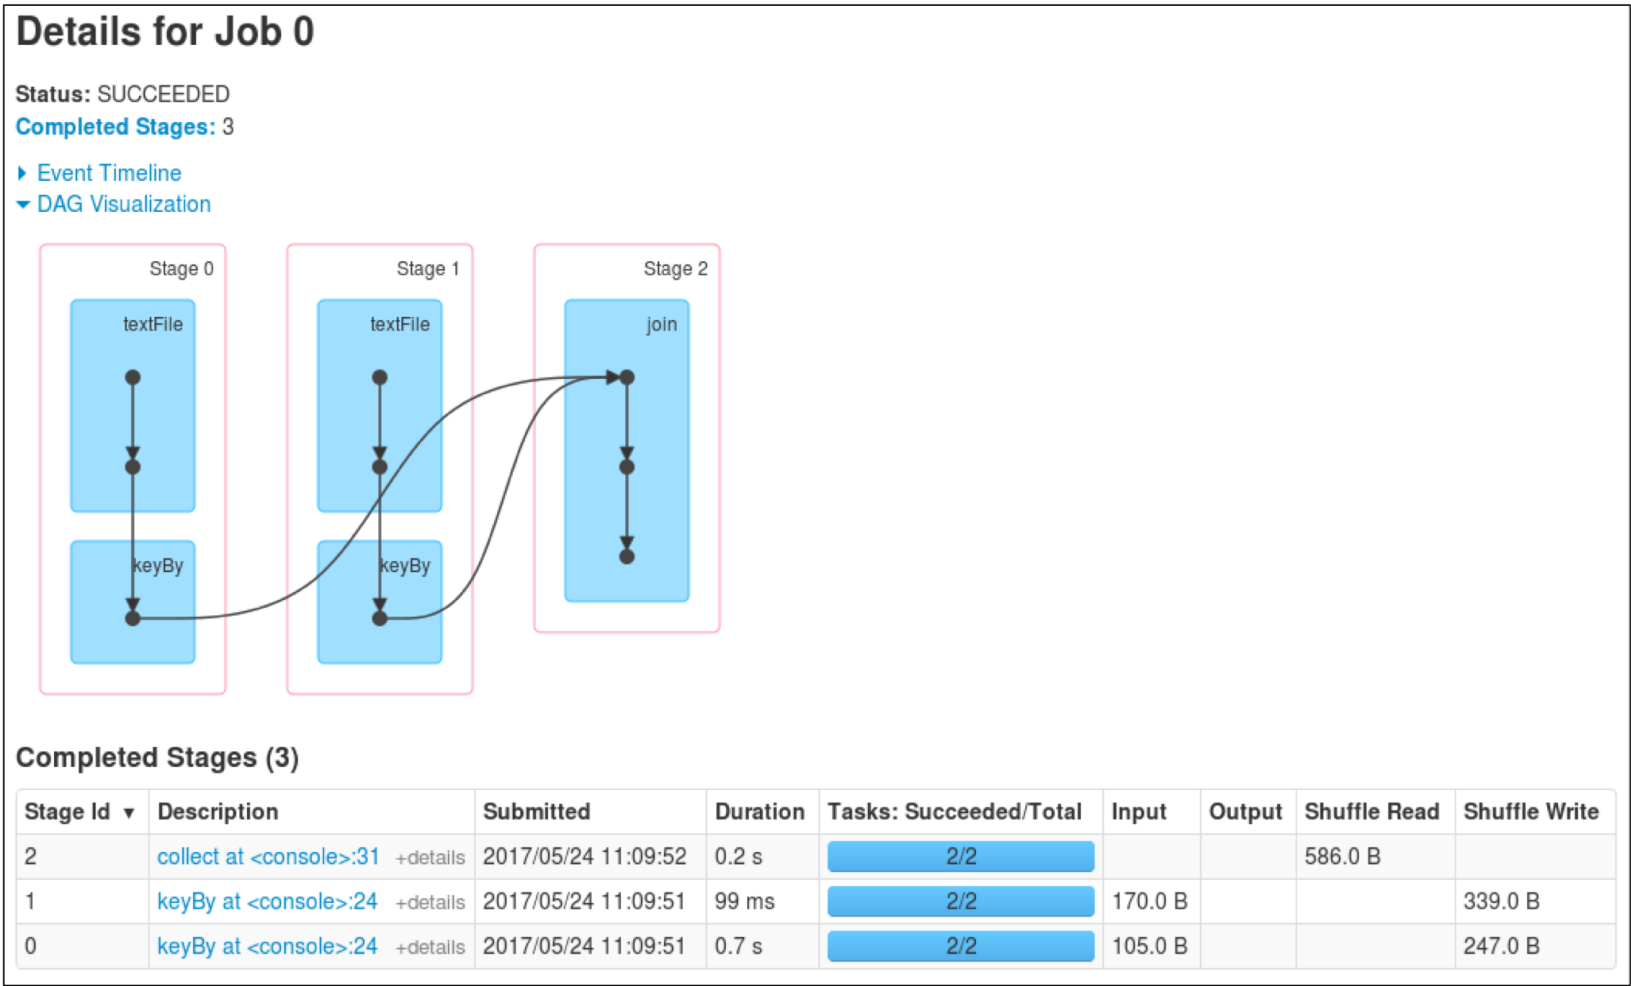
Spark Jobs (?)

User: training
Total Uptime: 55 s
Scheduling Mode: FIFO
Completed Jobs: 1
[Event Timeline](#)

Completed Jobs (1)

| Job Id ▼ | Description | Submitted | Duration | Stages: Succeeded/Total | Tasks (for all stages): Succeeded/Total |
|-------------|---|------------------------|----------|----------------------------|--|
| 0 | collect at <console>:31 | 2017/05/24 11:09:51 | 1 s | 3/3 | 6/6 |

Example: RDD Execution Plan (3)



Essential Points

- **Spark partitions split data across different executors in an application**
- **Executors execute query tasks that process the data in their partitions**
- **Narrow operations like `map` and `filter` are pipelined within a single stage**
 - Wide operations like `groupByKey` and `join` shuffle and repartition data between stages
- **Jobs consist of a sequence of stages triggered by a single action**
- **Jobs execute according to execution plans**
 - Core Spark creates RDD execution plans based on RDD lineages
 - Catalyst builds optimized query execution plans
- **You can explore how Spark executes queries in the Spark Application UI**

Hands-On Exercise: Jobs Monitoring : Using Web UI.

- **In this exercise, you will explore how Spark plans and executes RDD and DataFrame/Dataset queries**
 - Please refer to the Hands-On Exercise Manual for instructions